

— AV AI DISPATCH • RESEARCH REPORT

# The AV AI Production Index

A quarterly assessment of what AI can and can't do in audio/video production. No hype. Just what works, what's close, and what's still a demo.

Q1 2026 • INAUGURAL EDITION • AUDIO/VIDEO SERVICES

## 40 Years Later, the Same Lesson

*In 1986, a headquarters staff member told a writer, "It's sure great that you have MEPS. Now writing must be easy for you." The writer had to explain that people, not machines, do the writing. Forty years later, the technology is different, but the punchline is the same.*

**70–90%**

AI handles this much

**10–30%**

Still needs a human

**100%**

Needs QA before shipping

That last 10–30% is where quality lives. Every category in this report follows the same pattern.

# Where Things Stand **Right Now**

## ✓ WHAT YOU CAN DO

- Generate **5–15 seconds** of usable video footage for B-roll and concept work
- Produce **broadcast-quality background music** indistinguishable from stock libraries
- **Clone voices** with near-human fidelity from as little as 15 seconds of audio
- Cut **post-production time in half** with AI cleanup, captioning, and masking

## ✗ WHAT YOU CAN'T DO

- Produce a coherent **60-second narrative** scene
- Replace a **composer** for distinctive, memorable music
- Ship **any AI output** without a human quality pass

# Video Generation

The most visible category — and the one with the widest gap between demos and production reality.

# The Field Right Now

| TOOL           | RESOLUTION   | DURATION             | CONSISTENCY       | PRICE          |
|----------------|--------------|----------------------|-------------------|----------------|
| Runway Gen-4.5 | 4K           | Up to 60s            | Good              | \$12-\$76/mo   |
| Kling 3.0      | 4K at 60fps  | 15s                  | Good (multi-shot) | \$6-\$48/mo    |
| Seedance 2.0   | 2K           | 15s                  | Good (multi-ref)  | TBD            |
| Google Veo 3.1 | Up to 4K     | 4-8s                 | Moderate          | Vertex pricing |
| Sora 2         | 1080p        | 20s                  | Weak              | \$20-\$200/mo  |
| Pika 2.5       | 1080p        | 10s (25s Pikaframes) | Weak              | \$8-\$76/mo    |
| Luma Ray 3.14  | 1080p native | 5-10s (ext. 30s)     | Moderate          | Enterprise     |

# Readiness Breakdown

## READY

- **Short clips (5–15s)** for B-roll and supplemental content
- **Image-to-video** — consistently beats text-to-video
- **Cinematic style & camera movement**
- **Pre-vis & storyboarding** — ~40% time savings
- **Lip-sync** in controlled settings

## CLOSE

- **Character consistency** — drifts past 3–4 shots
- **Extended duration** — coherence fails past ~20s
- **Physics** — rigid OK, fluid/cloth fail

## NOT YET

- **Text in video** — all tools fail
- **Hands/fingers** — persists across all models
- **30s+ narrative coherence**

**Takeaway:** Use now for pre-vis and B-roll. Don't ship AI-only video for precise human representation.

# Editing & Post-Production

Where AI delivers the most immediate, measurable time savings today.

# Readiness Breakdown

## READY

- **Transcript-based editing (Descript)** — 60–70% time savings for spoken-word
- **Object masking (Premiere Pro)** — hover, click, tracked mask in seconds
- **AI audio cleanup (Descript)** — lifts bad audio to usable
- **Auto-captioning** — 92–95% accuracy

## CLOSE

- **Generative Extend** — subtle extensions only
- **AI dubbing** — strong for European languages, weak for tonal

## NOT YET

- **Creative editorial judgment** — pacing, arc, timing
- **Brand-consistent editing at scale**
- **Complex narrative editing**

**Takeaway:** The strongest ROI in this entire report. Transcript-based editing alone justifies the tool cost.



# VFX & Compositing

AI rotoscoping is the single most production-ready AI capability in all of VFX.

# Visual Effects & Animation

## VFX & Compositing

### READY

- **AI rotoscoping** — Premiere Object Mask, Slapshot (8K, 400 shots/day)
- **CG compositing (Wonder Studio)** — 80–90% automated, from \$19.99/mo
- **Simple object removal** — wires, boom mics

### NOT YET

- Fully automated VFX shots
- Real-time AI compositing

## Animation

### READY

- **Cascadeur** — AutoPosing + AutoPhysics, 2–3× faster blocking
- **Mocap cleanup** — Rokoko + Cascadeur, minutes not days
- **Single-camera mocap** — webcam previz

### NOT YET

- Fully automated character animation
- AI lip-sync for 3D characters

**Takeaway:** Rotoscoping and mocap cleanup are deploy-today tools. Final artistry still needs humans.

# Audio: Voice, Music & Post

The category where AI most consistently passes for human output.

# Voice Synthesis & Text-to-Speech

READY

- **Short-to-medium narration (<10 min)** — indistinguishable from human for explainer content
- **Voice cloning** — ElevenLabs (~30 min audio), Fish Audio S2 (15 seconds)
- **Word-level emotion control (Fish Audio S2)** — tag words with [whispered] , [voice breaking]

CLOSE

- Long-form narration (30+ min) — prosody drift accumulates
- Complex emotional arcs — irony, sarcasm, building tension

## Pricing Snapshot

| TIER           | COST                    |
|----------------|-------------------------|
| ElevenLabs     | From ~\$0.12/1K chars   |
| Fish Audio S2  | Free (self-host) or API |
| Azure / Google | ~\$0.016/1K chars       |

**Takeaway:** AI TTS is production-ready for narration under 10 minutes. Premium dramatic performance still needs a human.

# Music, SFX & Audio Post

## Music Generation

- **READY** **Background music & underscore** — Suno V5, 48kHz, 7M tracks/day
- **CLOSE** Full songs — vocals drift, structures collapse
- **NOT YET** Distinctive compositional voice

## Sound Design & SFX

- **READY** Simple ambience, UI sounds, atmospheres
- **NOT YET** Foley, cinematic sound design, anything that ships

## Audio Post-Production

The most production-ready AI audio category.

- **Noise removal (iZotope RX 11)** — industry gold standard
- **Speech enhancement (Adobe Podcast)** — free (1hr/day)
- **10-stem separation (LALAL.AI)** — 14.8M hours processed
- **Transcript editing (Descript)** — one-click cleanup

**Takeaway:** Audio post saves the most hours per dollar. iZotope RX 11 and Descript belong in every toolkit.

# Dubbing, Platforms & Agents

The pipeline is 4-5 tools stitched together. True end-to-end doesn't exist yet.

# The Pipeline Today

## Dubbing & Localization

- **READY** **AI dubbing + human QA** for informational content (Deepdub, Papercup, ElevenLabs)
- **CLOSE** Entertainment dubbing — sounds noticeably AI
- **NOT YET** Fully automated dubbing, singing, premium entertainment

## AI Agents

- **Make.com / n8n** — ~50% time reduction for trigger-based workflows
- No agent is ready for **autonomous production**

## End-to-End Platforms

| PLATFORM         | READY?             |
|------------------|--------------------|
| ElevenLabs Flows | Alpha              |
| Luma Agents      | Just launched      |
| Synthesia        | Yes — talking-head |
| HeyGen           | Yes — talking-head |

**Realistic pipeline:** Script (LLM) → Video (Synthesia/Runway) → Audio (ElevenLabs) → Edit (NLE/Descript) → Deliver

# AI Coding Agents & Dev Harnesses

Two layers: IDE harnesses where a developer steers, and autonomous agents that work independently and return finished code.



# IDE Harnesses — Developer in the Loop

| HARNESS            | KEY CAPABILITY                                                                | PRICE         |
|--------------------|-------------------------------------------------------------------------------|---------------|
| Cursor             | Agent mode, 8 parallel agents, cloud agents from browser/phone/Slack          | \$20–\$200/mo |
| Windsurf           | Cascade writes ~90% of code; "Memories" learn project conventions; Turbo Mode | \$15–\$200/mo |
| Google Antigravity | Agent-first: plan → execute → test → verify across editor, terminal, browser  | Free preview  |

These are not autocomplete tools — they are agentic environments where you describe intent and the AI plans, writes, tests, and debugs. Over half the Fortune 500 use Cursor alone.

## Autonomous Agents — No IDE Required

| AGENT         | HOW IT WORKS                                                               | ACCESS                      |
|---------------|----------------------------------------------------------------------------|-----------------------------|
| OpenAI Codex  | Clones repo, writes code, runs tests, returns PR. GPT-5.3–Codex            | ChatGPT Plus (\$20/mo+)     |
| Claude Code   | Terminal agent: reads codebase, plans, executes multi-step implementations | Claude Max (\$100–\$200/mo) |
| Jules 3.0     | Async agent: assign from browser/Slack, returns PRs. API for CI/CD         | Free preview                |
| Claude Cowork | Knowledge work: file management, spreadsheets, reports — not code-specific | Claude Max (\$100–\$200/mo) |

# Use Cases for AV Teams

- **Internal tool development** — Asset management, review workflows, dashboards, metadata tagging. One developer + agent ≈ a team of three
  - **Web development & publishing** — Sites, CMS interfaces, distribution platforms. Agent handles scaffolding, responsive layouts, API integrations, deploys
  - **Automation infrastructure** — Ingest pipelines, transcription routing, QC checks. Codex/Jules work async on GitHub issues, return working PRs
  - **Production operations (Cowork)** — Organize project files, compile reports, generate formatted spreadsheets, manage docs
- NOT YET
- **Complex architecture decisions** — Deep domain knowledge, creative problem-solving
  - **Production-critical code without review** — Error rate on complex tasks still too high

**Takeaway:** IDE harnesses (Cursor, Windsurf) are the fastest path to internal tools. Autonomous agents (Codex, Jules) handle well-scoped tasks async. Neither replaces a developer — they replace the repetitive 70%.

# Transcription & Captioning

The most mature AI category in production.

| PROVIDER                   | ENGLISH ACCURACY | PRICE/MIN | BEST FOR                          |
|----------------------------|------------------|-----------|-----------------------------------|
| AssemblyAI Universal-3 Pro | 94.1%            | \$0.0035  | Audio intelligence, PII redaction |
| OpenAI Whisper             | ~92%             | \$0.006   | Budget batch, 100+ languages      |
| Deepgram Nova-3            | ~92%             | \$0.0043  | Real-time (<300ms latency)        |

**Weak point:** Speaker diarization — 85–90% in good conditions, 75–80% with noise or 5+ speakers.

**Takeaway:** Transcription is solved for most use cases. Budget under \$1/hour for enterprise-grade accuracy.



# The Cheat Sheet

What to adopt, what to watch, and what to skip.

ADOPT NOW

# Clear ROI

|                          |                                      |
|--------------------------|--------------------------------------|
| Concept / pre-vis        | Runway Gen-4.5 · Kling 3.0           |
| Rotoscoping & masking    | Premiere Object Mask · Slapshot      |
| Audio cleanup            | iZotope RX 11 · Adobe Podcast (free) |
| Transcript-based editing | Descript                             |
| Background music         | Suno V5                              |

|                              |                            |
|------------------------------|----------------------------|
| Voice synthesis (short-form) | ElevenLabs · Fish Audio S2 |
| Talking-head video           | Synthesia · HeyGen         |
| Animation blocking + mocap   | Cascadeur · Rokoko         |
| CG character compositing     | Wonder Studio              |
| Internal tool / web dev      | Cursor · Windsurf          |
| Async bug fixes & features   | Codex · Jules 3.0          |

WATCH CLOSELY

Agent-based workflows (Luma, ElevenLabs Flows) · Multi-shot narrative (Kling, Seedance) · AI dubbing at scale · Autonomous agents for complex dev

NOT YET

## Don't Count On

- AI-only long-form video generation (>30s coherent narrative)
- Automated editorial decision-making
- AI text rendering within generated video
- Unsupervised VFX compositing
- AI-only character animation for broadcast
- Fully autonomous production agents for creative work
- AI music as hero content
- AI sound effects for final delivery
- Autonomous coding agents for complex architecture

**Shelf life:** This assessment is useful for roughly 3–6 months. The Q2 2026 edition will follow.

— AV AI DISPATCH • RESEARCH REPORT

# The AV AI Production Index

Q1 2026 · Inaugural Edition

Full report available at [avaidispatch.com](https://avaidispatch.com)

AUDIO/VIDEO SERVICES — AI RESEARCH & DEVELOPMENT 2026